

# Hyperparameter specification experiment results

## Summary

In comparing the current and a new method for specifying hyperparameters to `covdepGE`, I found that although the current method produces greater ELBO, the new method consistently outperforms the current one in terms of accuracy. I include visualizations for the differences in metrics, as well as a section analyzing the distribution of the candidate hyperparameters proposed from `varbvs`.

## Experiment Overview

In this experiment, I compared two methods of hyperparameter specification across 500 trials. In each trial, the data was generated by setting a seed and then following the continuous data paradigm. That is, all of the individuals in the first and third interval of the covariate space had the same precision matrix ( $\Omega_1$  and  $\Omega_3$ ), while the precision matrix for individuals in the second interval continuously shifted from  $\Omega_1$  to  $\Omega_3$ .

The first method (M1) follows the methodology for hyperparameter specification that we have already been utilizing. For each variable, a grid of 20 slab variances  $\sigma_\beta^2$  are generated independent of the data and the other hyperparameters. The best  $\sigma_\beta^2$  is selected based on ELBO.  $\sigma^2$  is chosen by taking the mean across the candidates fit by `varbvs` and  $\pi$  is chosen by taking the mean across the grid generated by `varbvs`.

To demonstrate M1, below are the results for the first two variables from one trial. For each variable, the selected value of the hyperparameters are listed, followed by a value `hyperparameters` that lists each of the elements of the hyperparameter space for which a model was fit, along with the resulting ELBO.

```
out1 <- covdepGE(X, Z, var_min = 1e-3, var_max = 1, n_sigma = 20, CS = T)
```

```
##      |
```

```
out1$CAVI_details[1:2]
```

```
## [[1]]
## [[1]]$sigmasq
## [1] 0.7221278
##
## [[1]]$sigmabeta_sq
## [1] 0.1623777
##
## [[1]]$pi
## [1] 0.1398245
##
## [[1]]$ELBO
## [1] -14150.06
##
## [[1]]$converged_iter
```

```

## [1] 20
##
## [[1]]$hyperparameters
##      sigmasq sigmabetasq      pi      elbo
## [1,] 0.7221278 1.000000000 0.1398245 -14221.08
## [2,] 0.7221278 0.695192796 0.1398245 -14198.04
## [3,] 0.7221278 0.483293024 0.1398245 -14177.72
## [4,] 0.7221278 0.335981829 0.1398245 -14161.50
## [5,] 0.7221278 0.233572147 0.1398245 -14151.36
## [6,] 0.7221278 0.162377674 0.1398245 -14150.06
## [7,] 0.7221278 0.112883789 0.1398245 -14161.37
## [8,] 0.7221278 0.078475997 0.1398245 -14190.24
## [9,] 0.7221278 0.054555948 0.1398245 -14242.90
## [10,] 0.7221278 0.037926902 0.1398245 -14326.57
## [11,] 0.7221278 0.026366509 0.1398245 -14448.62
## [12,] 0.7221278 0.018329807 0.1398245 -14615.04
## [13,] 0.7221278 0.012742750 0.1398245 -14828.02
## [14,] 0.7221278 0.008858668 0.1398245 -15083.91
## [15,] 0.7221278 0.006158482 0.1398245 -15372.63
## [16,] 0.7221278 0.004281332 0.1398245 -15678.96
## [17,] 0.7221278 0.002976351 0.1398245 -15984.49
## [18,] 0.7221278 0.002069138 0.1398245 -16270.49
## [19,] 0.7221278 0.001438450 0.1398245 -16521.11
## [20,] 0.7221278 0.001000000 0.1398245 -16725.53
##
##
## [[2]]
## [[2]]$sigmasq
## [1] 0.5723876
##
## [[2]]$sigmabeta_sq
## [1] 0.483293
##
## [[2]]$pi
## [1] 0.1398245
##
## [[2]]$ELBO
## [1] -17084.37
##
## [[2]]$converged_iter
## [1] 33
##
## [[2]]$hyperparameters
##      sigmasq sigmabetasq      pi      elbo
## [1,] 0.5723876 1.000000000 0.1398245 -17113.63
## [2,] 0.5723876 0.695192796 0.1398245 -17093.11
## [3,] 0.5723876 0.483293024 0.1398245 -17084.37
## [4,] 0.5723876 0.335981829 0.1398245 -17092.66
## [5,] 0.5723876 0.233572147 0.1398245 -17125.30
## [6,] 0.5723876 0.162377674 0.1398245 -17192.39
## [7,] 0.5723876 0.112883789 0.1398245 -17307.53
## [8,] 0.5723876 0.078475997 0.1398245 -17488.48
## [9,] 0.5723876 0.054555948 0.1398245 -17757.30
## [10,] 0.5723876 0.037926902 0.1398245 -18139.46

```

```
## [11,] 0.5723876 0.026366509 0.1398245 -18661.02
## [12,] 0.5723876 0.018329807 0.1398245 -19343.11
## [13,] 0.5723876 0.012742750 0.1398245 -20193.69
## [14,] 0.5723876 0.008858668 0.1398245 -21198.65
## [15,] 0.5723876 0.006158482 0.1398245 -22316.71
## [16,] 0.5723876 0.004281332 0.1398245 -23483.29
## [17,] 0.5723876 0.002976351 0.1398245 -24624.05
## [18,] 0.5723876 0.002069138 0.1398245 -25671.98
## [19,] 0.5723876 0.001438450 0.1398245 -26580.43
## [20,] 0.5723876 0.001000000 0.1398245 -27327.71
```

The second method, M2, relies entirely on `varbvs` to generate the grid of hyperparameters by fitting  $\sigma^2$  and  $\sigma_\beta^2$  to the data for each value of  $\pi$  in a grid of length 20. Again, the results for the first two variables are displayed.

```
out2 <- covdepGE(X, Z)
```

```
## |
```

```
out2$CAVI_details[1:2]
```

```
## [[1]]
## [[1]]$sigmasq
## [1] 0.7214097
##
## [[1]]$sigmabeta_sq
## [1] 0.9839818
##
## [[1]]$pi
## [1] 0.1923954
##
## [[1]]$ELBO
## [1] -14212.6
##
## [[1]]$converged_iter
## [1] 18
##
## [[1]]$hyperparameters
##      sigmasq sigmabetasq      pi      elbo
## [1,] 0.7212941  0.9832076 0.20000000 -14212.90
## [2,] 0.7214097  0.9839818 0.19239536 -14212.60
## [3,] 0.7215196  0.9847206 0.18501301 -14212.61
## [4,] 0.7216240  0.9854256 0.17785154 -14212.90
## [5,] 0.7217233  0.9860982 0.17090914 -14213.47
## [6,] 0.7218178  0.9867401 0.16418362 -14214.32
## [7,] 0.7219076  0.9873527 0.15767243 -14215.41
## [8,] 0.7219930  0.9879372 0.15137269 -14216.76
## [9,] 0.7220744  0.9884950 0.14528125 -14218.34
## [10,] 0.7221518  0.9890272 0.13939466 -14220.16
## [11,] 0.7222255  0.9895351 0.13370928 -14222.19
## [12,] 0.7222956  0.9900197 0.12822124 -14224.43
## [13,] 0.7223624  0.9904824 0.12292649 -14226.88
```

```

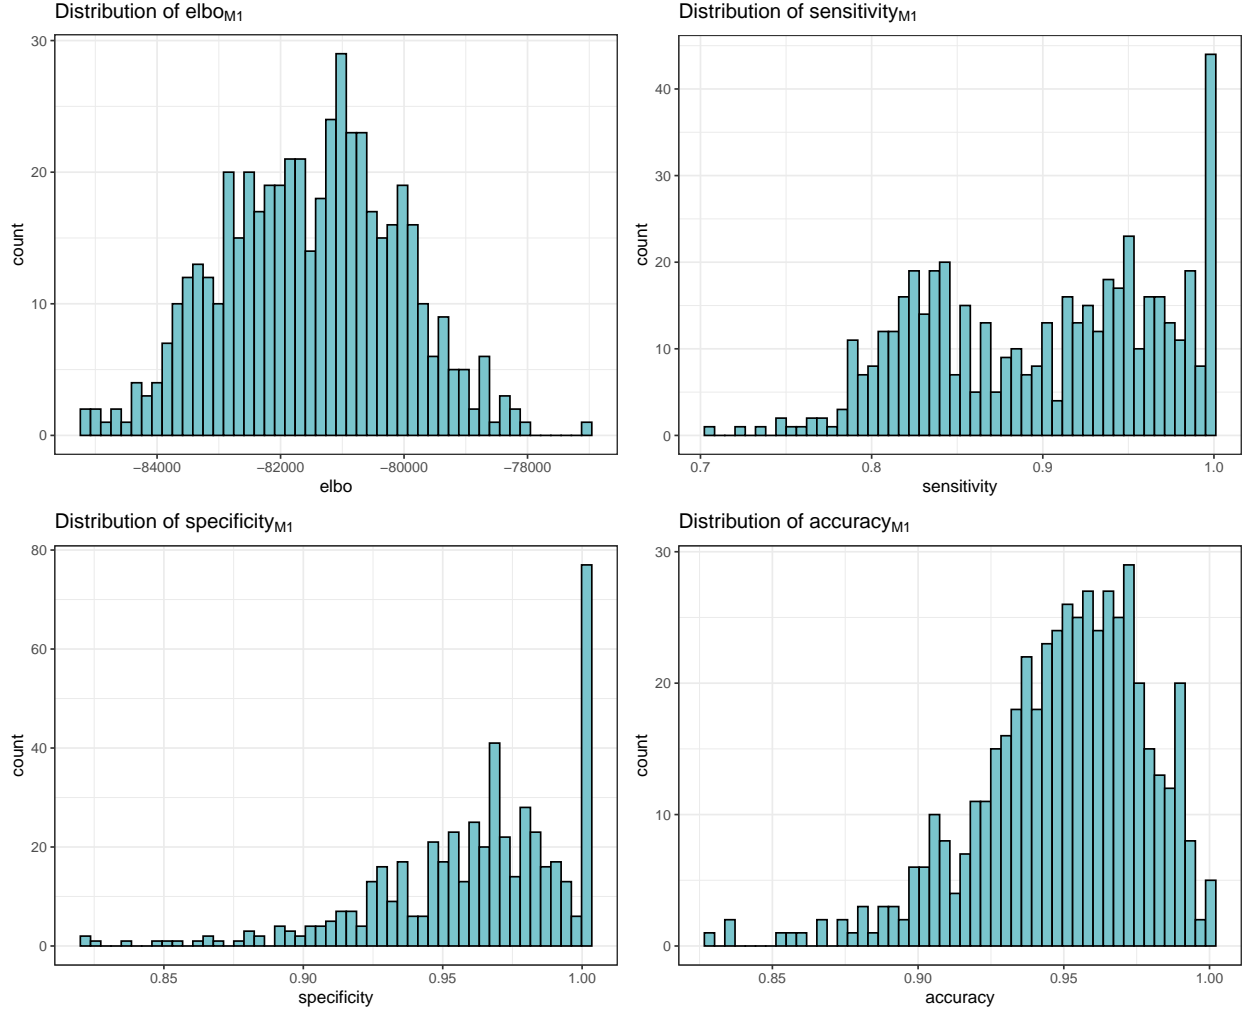
## [14,] 0.7224260 0.9909236 0.11782083 -14229.52
## [15,] 0.7224866 0.9913445 0.11289993 -14232.35
## [16,] 0.7225442 0.9917461 0.10815936 -14235.36
## [17,] 0.7225992 0.9921293 0.10359460 -14238.54
## [18,] 0.7226515 0.9924949 0.09920105 -14241.89
## [19,] 0.7227015 0.9928441 0.09497411 -14245.39
## [20,] 0.7227488 0.9931764 0.09090909 -14249.05
##
##
## [[2]]
## [[2]]$sigmasq
## [1] 0.5760302
##
## [[2]]$sigmabeta_sq
## [1] 0.9574463
##
## [[2]]$pi
## [1] 0.09090909
##
## [[2]]$ELBO
## [1] -17100.79
##
## [[2]]$converged_iter
## [1] 32
##
## [[2]]$hyperparameters
##      sigmasq sigmabetasq      pi      elbo
## [1,] 0.5692560 0.9482048 0.20000000 -17129.56
## [2,] 0.5695278 0.9486262 0.19239536 -17128.84
## [3,] 0.5698092 0.9490533 0.18501301 -17128.06
## [4,] 0.5701031 0.9494897 0.17785154 -17127.12
## [5,] 0.5704050 0.9499300 0.17090914 -17126.15
## [6,] 0.5707163 0.9503774 0.16418362 -17125.11
## [7,] 0.5710374 0.9508319 0.15767243 -17123.97
## [8,] 0.5713681 0.9512937 0.15137269 -17122.74
## [9,] 0.5717084 0.9517632 0.14528125 -17121.41
## [10,] 0.5720582 0.9522406 0.13939466 -17119.99
## [11,] 0.5724174 0.9527258 0.13370928 -17118.46
## [12,] 0.5727858 0.9532191 0.12822124 -17116.84
## [13,] 0.5731632 0.9537204 0.12292649 -17115.12
## [14,] 0.5735493 0.9542296 0.11782083 -17113.31
## [15,] 0.5739437 0.9547468 0.11289993 -17111.41
## [16,] 0.5743462 0.9552717 0.10815936 -17109.44
## [17,] 0.5747562 0.9558040 0.10359460 -17107.38
## [18,] 0.5751732 0.9563435 0.09920105 -17105.27
## [19,] 0.5755969 0.9568898 0.09497411 -17103.10
## [20,] 0.5760302 0.9574463 0.09090909 -17100.79


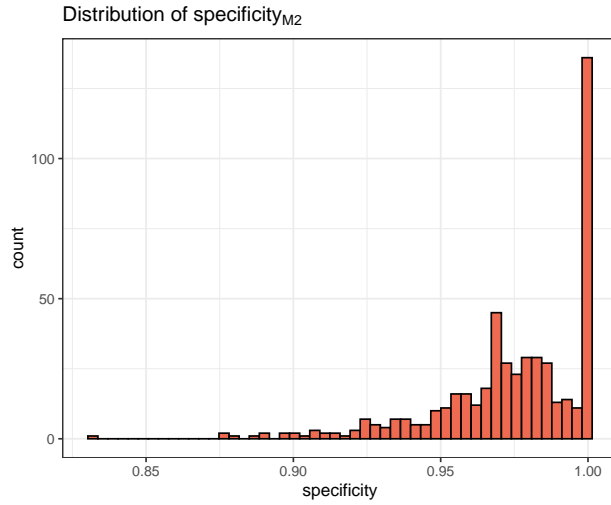
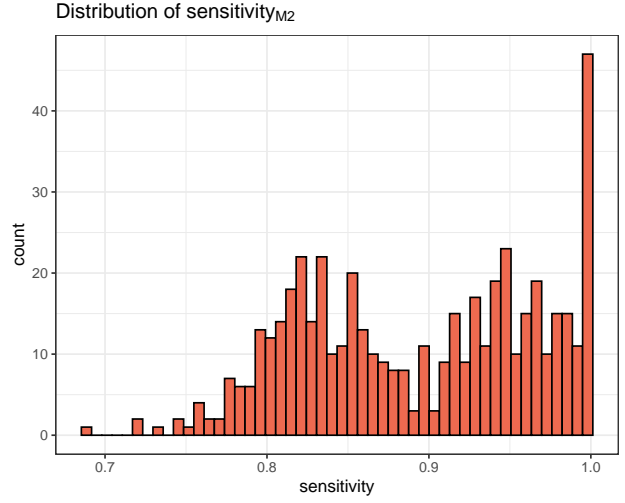
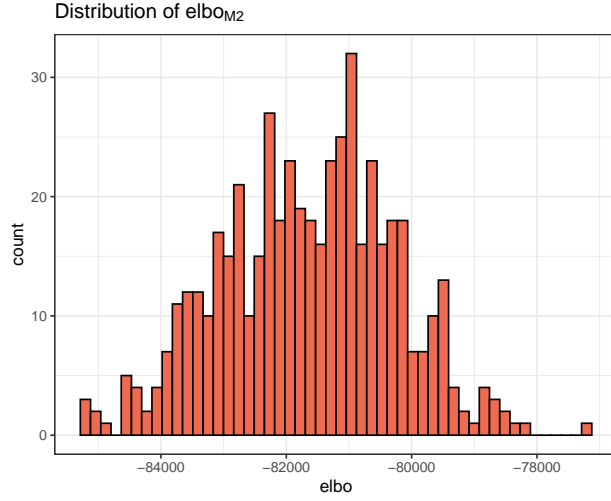
```

## Results: ELBO and accuracy

First, as a frame of reference for scale, the distribution of the ELBO, sensitivity, specificity, and accuracy for both M1 and M2 is visualized.

*Note: One of the ELBO for  $M2$  was removed, as it was on the order of  $1 \times 10^{12}$ . However, all other metrics for this model were within 3% of that of  $M1$ , and so these were not dropped*

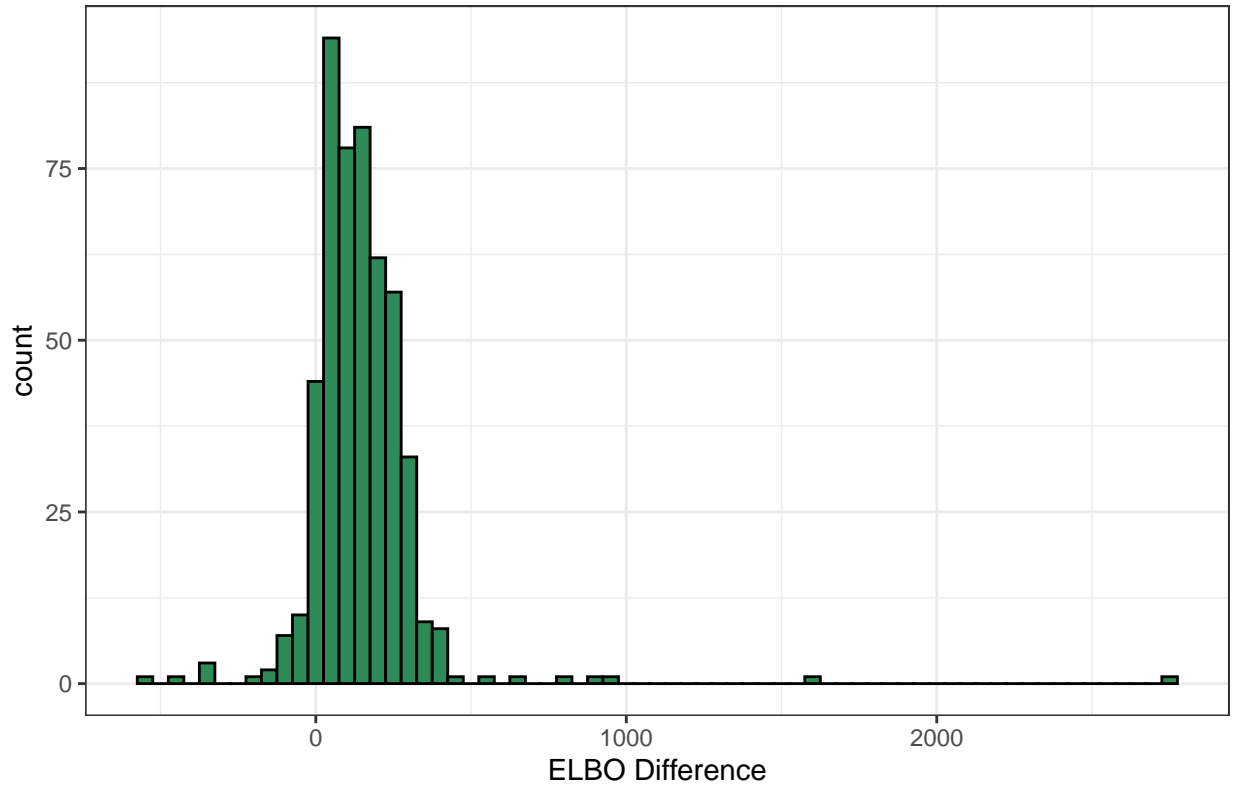




Visualized next is the pairwise difference between the ELBO of M1 and M2. The ELBO of M1 was more frequently greater than that of M2

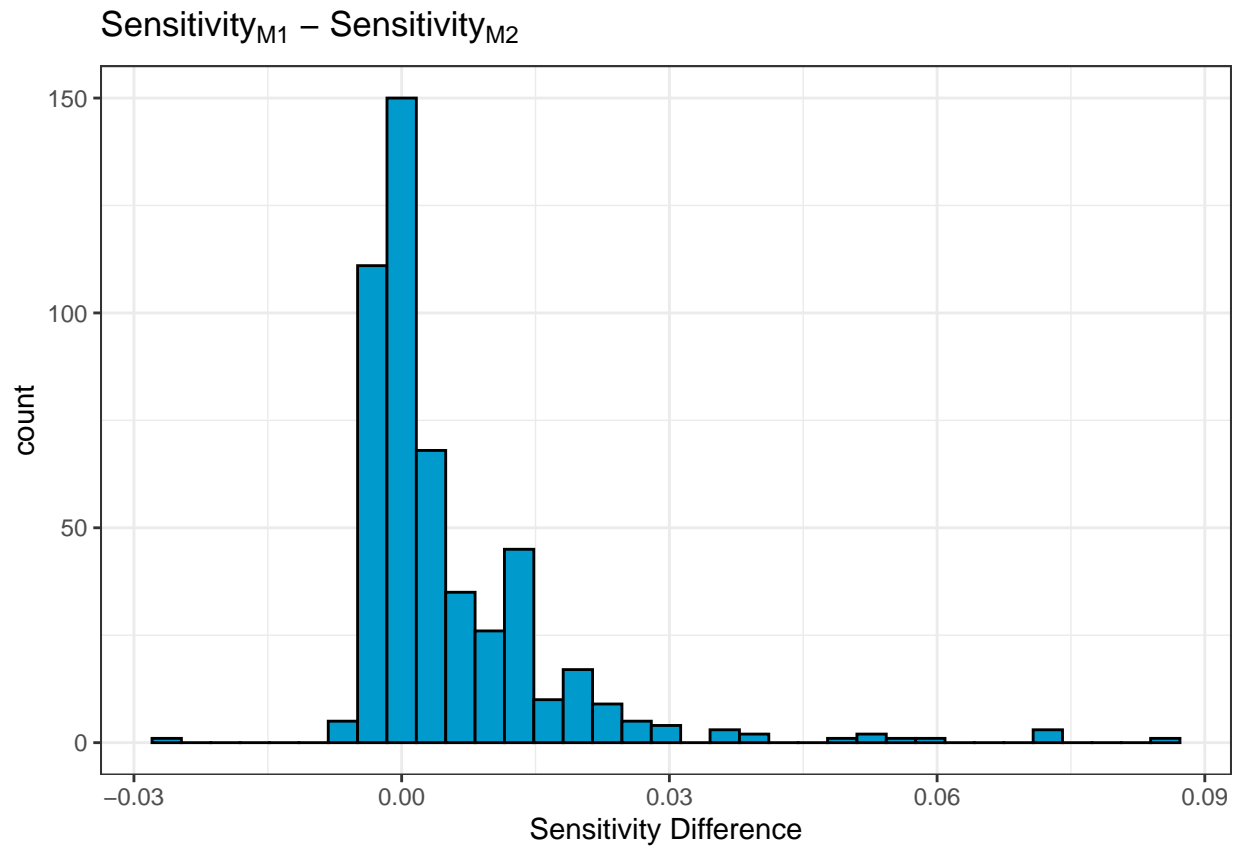
*Note: the large ELBO for M2 was again removed*

Distribution of  $ELBO_{M1} - ELBO_{M2}$



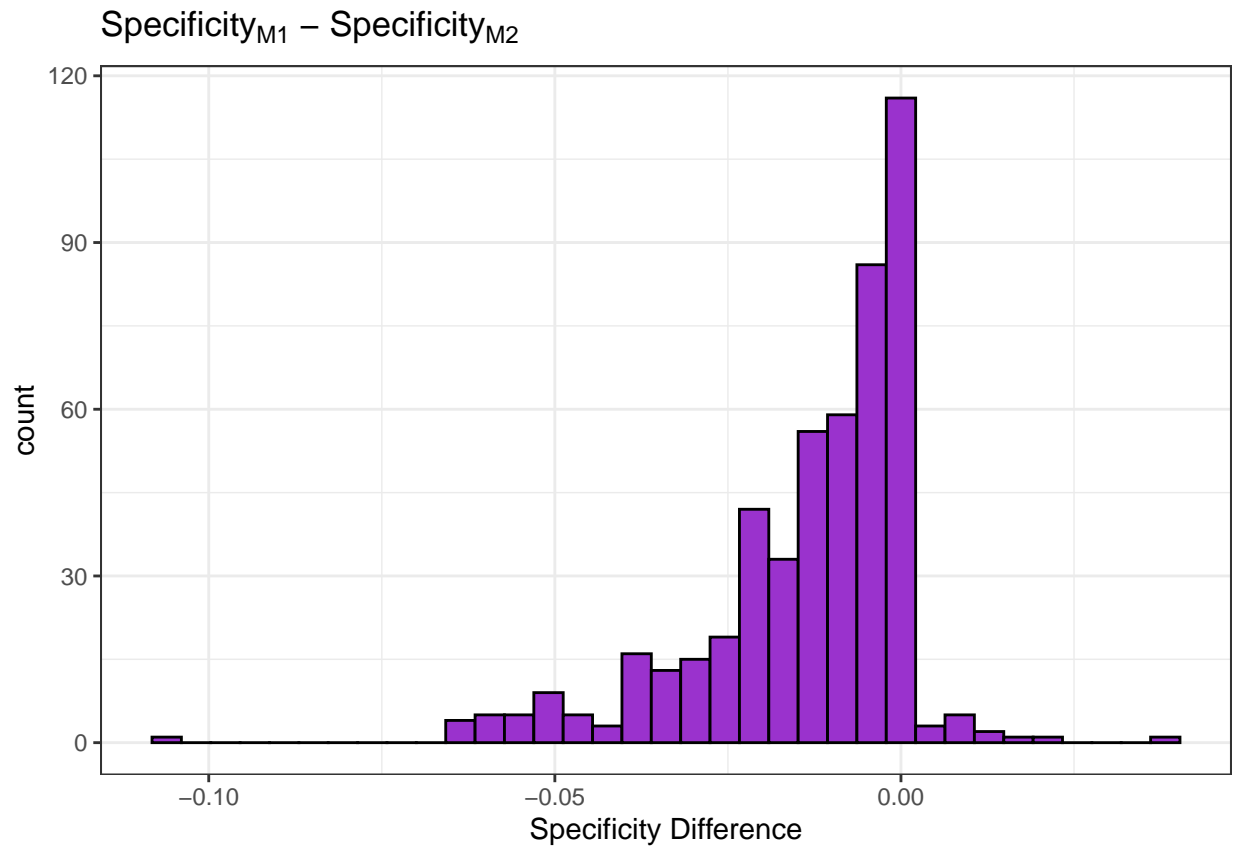
```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -543.2   58.3   134.5   146.4   217.3   2764.3
```

Visualized next is the pairwise differences in sensitivity, specificity, and accuracy. Although the rate at which edges were correctly detected (sensitivity) was roughly equivalent between the methods, M2 outperformed M1 in non-edge detection rate (specificity), which resulted in the superior accuracy of M2.

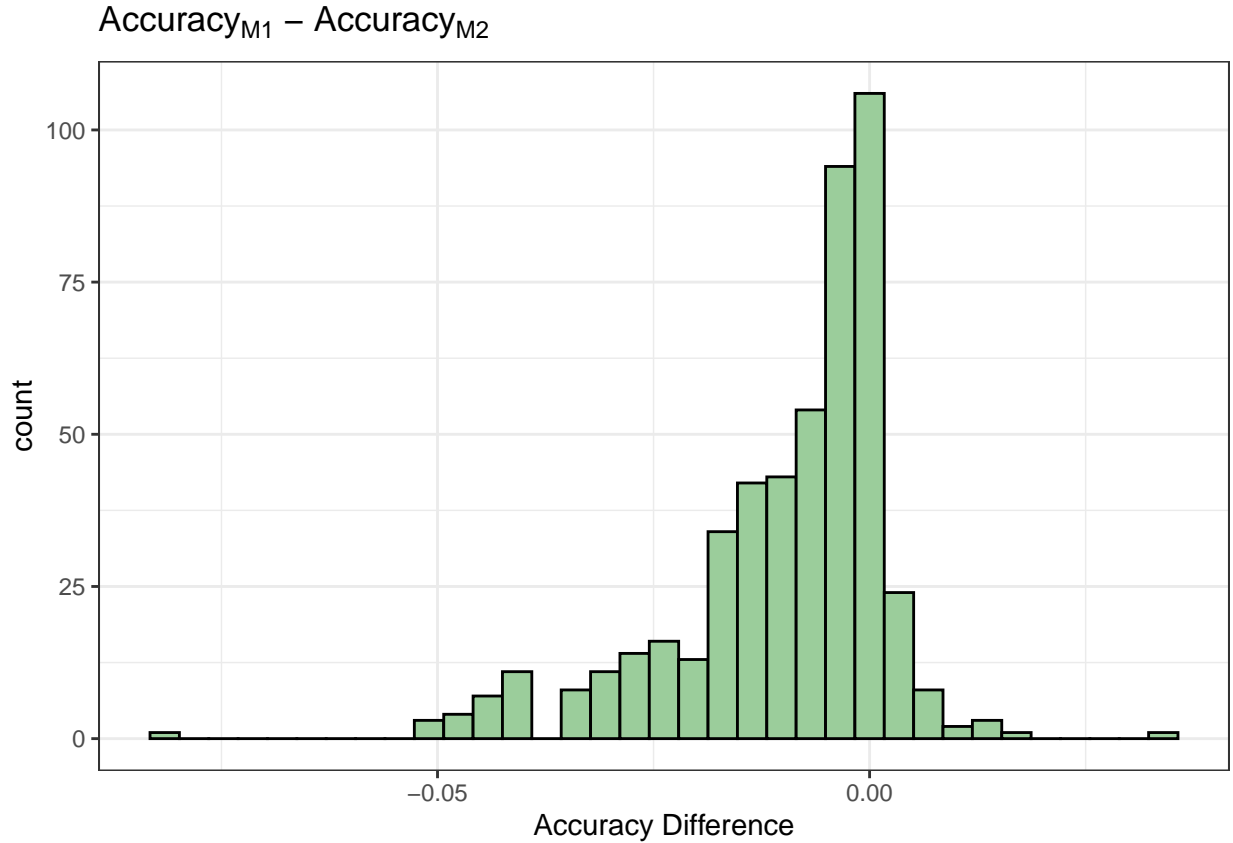


##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	-0.026190	0.000000	0.000000	0.005333	0.009524	0.085714





```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## -0.104372 -0.020219 -0.008743 -0.013143 -0.001639  0.039891
```



```
##      Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
## -0.0813333 -0.0151111 -0.0057778 -0.0096942 -0.0004444  0.0342222
```

## Distribution of Candidates

The  $\pi$  candidates generated by `varbvs` result from a grid of log-odds spanning from  $-\log_{10}(p)$  to  $-\log_{10}(1)$ :

```
p <- 4
log.odds <- seq(-log10(p), -1, length.out = 20)
(pi <- 1 / (1 + 10^(-log.odds)))

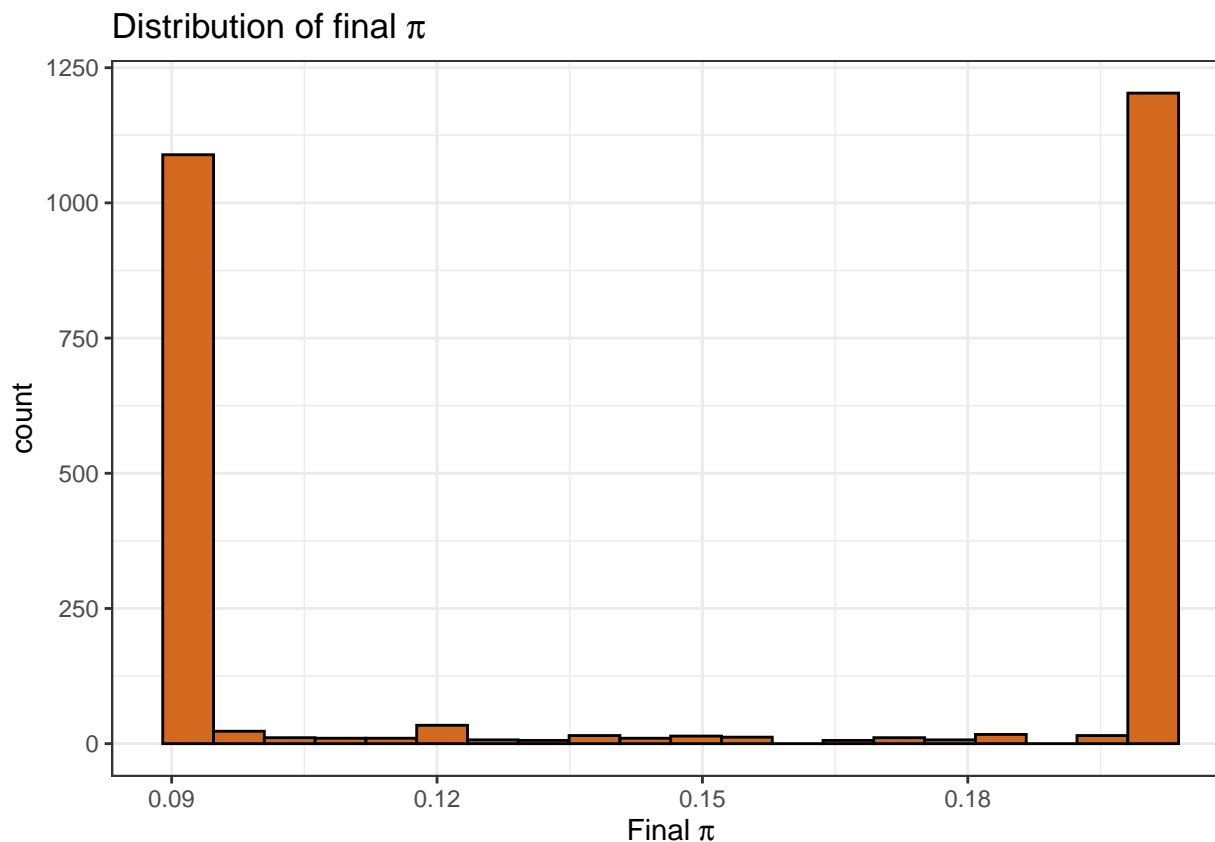
## [1] 0.20000000 0.19239536 0.18501301 0.17785154 0.17090914 0.16418362
## [7] 0.15767243 0.15137269 0.14528125 0.13939466 0.13370928 0.12822124
## [13] 0.12292649 0.11782083 0.11289993 0.10815936 0.10359460 0.09920105
## [19] 0.09497411 0.09090909
```

```
mean(pi)
```

```
## [1] 0.1398245
```

Since M1 uses the mean of these candidates, all of the hyperparameter candidates for M1 use  $\pi \approx 0.14$ , as can be seen from the demonstration of M1 in the experiment overview section. However, M2 fits a model for each of the candidates in the  $\pi$  grid.

The following is a visualization of the final values of  $\pi$  chosen by the M2 models. Since each variable can have its own value of  $\pi$ , and there are 5 variables per model, this distribution contains 2,500 values of  $\pi$ . This visualization shows that ELBO favors values of  $\pi$  on the extremes of the distribution.



I proposed M2 with the reasoning that since it fits  $\sigma^2$  and  $\sigma_\beta^2$  to the data for each  $\pi$  in the grid, it would cover a greater portion of the hyperparameter space than M1. It is important to note that after analyzing the spread of  $\sigma^2$  and  $\sigma_\beta^2$  fit to each variable by `varbvs`, it became apparent that the candidates were relatively similar to one another.

For reference to the magnitude of these candidates, here are summary statistics for the  $\sigma^2$  candidates fit to the data by `varbvs`:

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.3268 0.4804 0.5214 0.5327 0.5722 0.9960
```

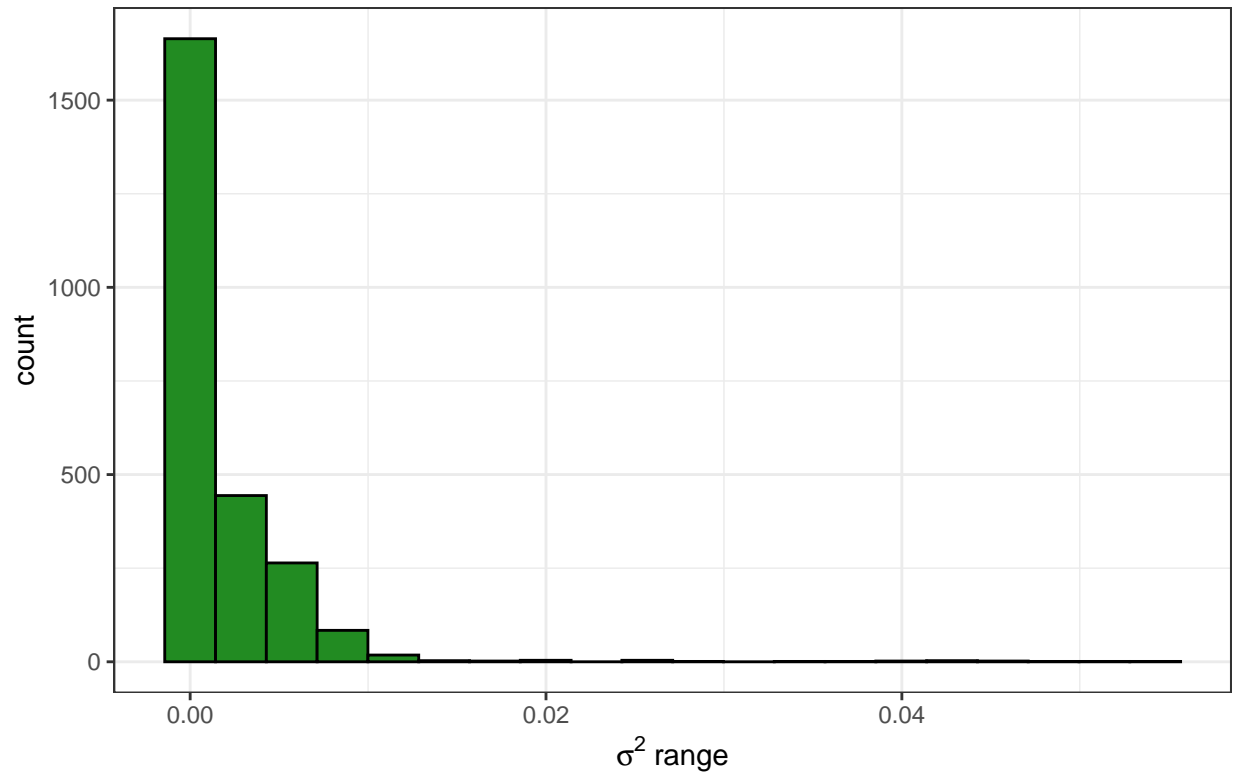
And the  $\sigma_\beta^2$  candidates:

```
summary(as.numeric(res$hyperparameters$sigmabetasq_cands))
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.8692 0.9401 0.9636 0.9630 0.9944 1.0039
```


The following visualizations are the distribution of the ranges of these candidates proposed by `varbvs`. For each variable in a model, `varbvs` fits  $\sigma^2$  and  $\sigma_\beta^2$  to the data for each of the 20  $\pi$  values in the grid. I defined the range as the difference between the largest and smallest hyperparameter fit to the data; since there were 500 models, and 5 variables per model, both of these distributions represent 2,500 ranges.

Distribution of ranges of  $\sigma^2$



##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	6.000e-08	1.578e-04	5.914e-04	1.994e-03	2.470e-03	5.424e-02

Distribution of ranges of  $\sigma_{\beta}^2$



```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.001105 0.002679 0.003999 0.005806 0.007253 0.079401
```